

# **FINAL PROJECT REPORT**

**MEDICAL INSURANCE EXPENSE**

**HA HAI VU**

## Data gathering and intergration:

For this project, I pick dataset *insurance.csv* from the book “*Machine Learning with R*” by Brett Lantz. I have discovered this data from GitHub, the link to the documentation is [here](#).

This dataset includes 1,338 examples of beneficiaries currently enrolled in the insurance plan, with 7 features indicating characteristics of the patient as well as the total medical expenses charged to the plan for the calendar year. The features are:

- age: age of the primary beneficiary
- sex: insurance contractor gender, female, male
- bmi: body mass index, providing an understanding of body, weights that are relatively high or low relative to height, objective index of body weight ( $\frac{kg}{m^2}$ ) using the ratio of height to weight, ideally 18.5 to 24.9.
- children: Number of children covered by health insurance / Number of dependents
- smoker: this is yes/no depending on whether the insured regularly smokes tobacco.
- region: the beneficiary's residential area in the US, northeast, southeast, southwest, northwest
- charges: individual medical costs billed by health insurance

I have decided to use this dataset to predict insurance costs based on the provided content.

```
> setwd("D:/E - Working/@Master - Data Science/Study/7. DSC 441 - Fundamentals of Data Science/Week 10 - 03.07.2022/HW5")
> insurance = read.csv("insurance.csv")
> head(insurance)
  age  sex  bmi children smoker  region  charges
1  19 female 27.900      0   yes southwest 16884.924
2  18  male 33.770      1    no southeast 1725.552
3  28  male 33.000      3    no southeast 4449.462
4  33  male 22.705      0    no northwest 21984.471
5  32  male 28.880      0    no northwest 3866.855
6  31 female 25.740      0    no southeast 3756.622
```

## Data exploration:

```
> str(insurance)
'data.frame': 1338 obs. of 7 variables:
 $ age      : int  19 18 28 33 32 31 46 37 37 60 ...
 $ sex      : chr   "female" "male" "male" "male" ...
 $ bmi      : num   27.9 33.8 33 22.7 28.9 ...
 $ children: int    0 1 3 0 0 0 1 3 2 0 ...
 $ smoker   : chr   "yes" "no" "no" "no" ...
 $ region   : chr   "southwest" "southeast" "southeast" "northwest" ...
 $ charges  : num  16885 1726 4449 21984 3867 ...
```

```
> summary(insurance)
```

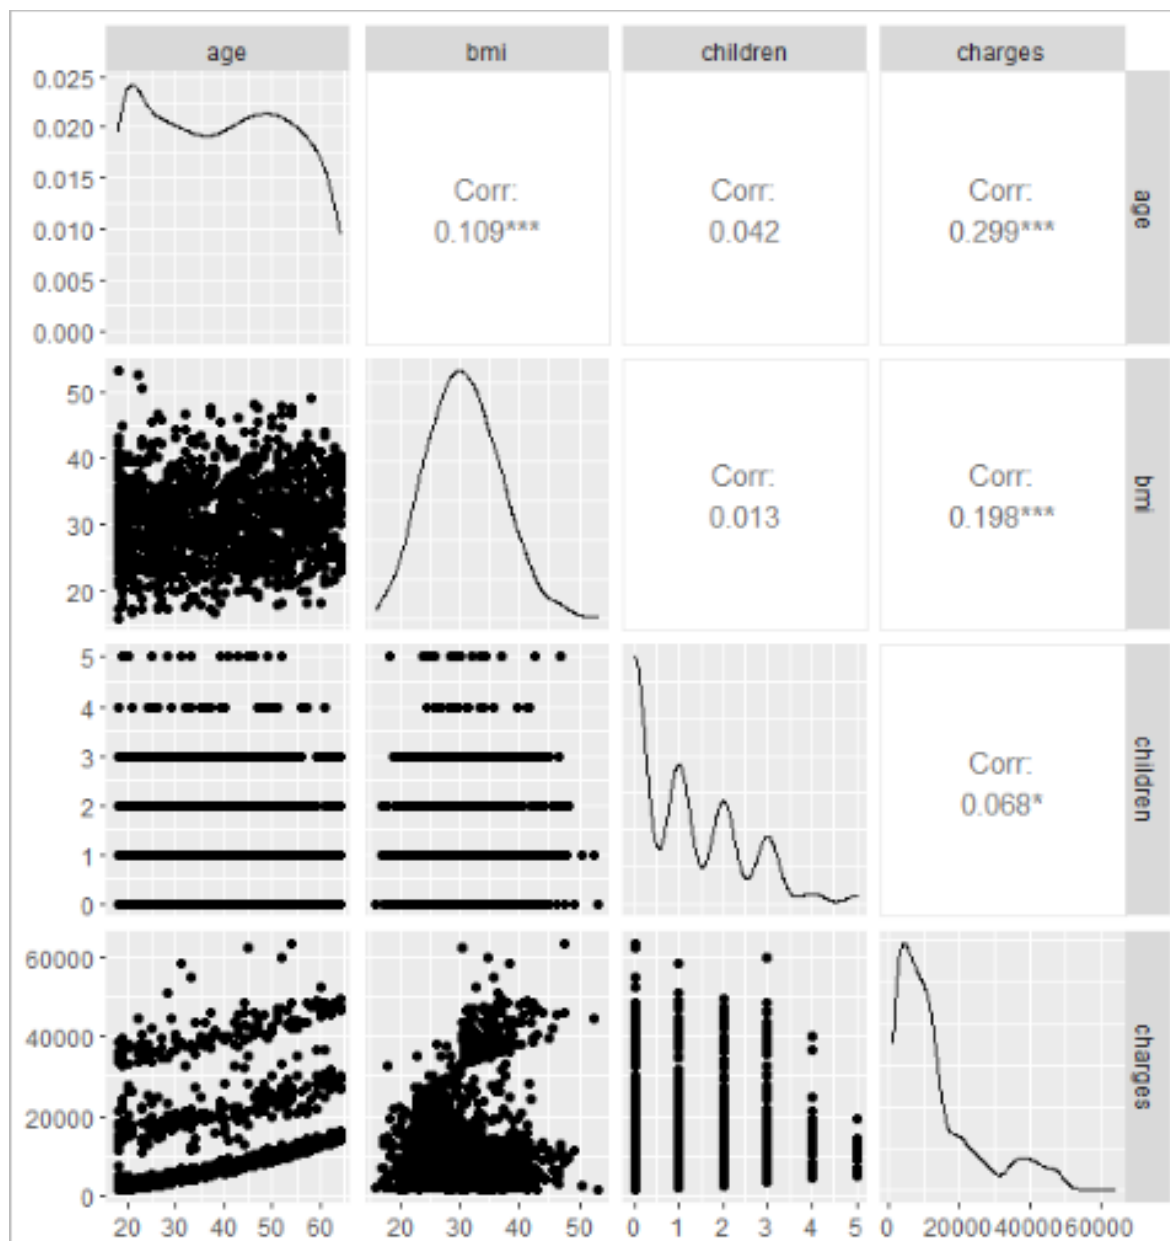
age		sex	bmi		children	smoker	region	charges
Min.	:18.00	Length:1338	Min.	:15.96	Min. :0.000	Length:1338	Length:1338	Min. : 1122
1st Qu.	:27.00	Class :character	1st Qu.	:26.30	1st Qu.:0.000	Class :character	Class :character	1st Qu.: 4740
Median	:39.00	Mode :character	Median	:30.40	Median :1.000	Mode :character	Mode :character	Median : 9382
Mean	:39.21		Mean	:30.66	Mean :1.095			Mean :13270
3rd Qu.	:51.00		3rd Qu.	:34.69	3rd Qu.:2.000			3rd Qu.:16640
Max.	:64.00		Max.	:53.13	Max. :5.000			Max. :63770

We can see that there are no missing values in this dataset. In this data, the types of variables included are categorical (sex, smoker, region) and numeric (age, bmi, children and charges).

Now, let's perform some more exploratory tasks and find existing relationships.

First, I'll create a scatterplot matrix of the numerical variables:

```
insurance %>% select(age, bmi, children, charges) %>% ggpairs()
```




We can see a few things worth from this graph: Firstly, while there are outliers, we can see graphically that charges generally increase with age. The Pearson correlation coefficient also shows a positive correlation between charges and age. Secondly, the trend between BMI and charges isn't that clear graphically, but we can see that there is definitely an increase in charges for some individuals after hitting the age of 30. The Pearson correlation coefficient also proves a somewhat positive correlation. This will be interesting to discover later based on sex and smoker status. Thirdly, we note the somewhat positive correlation number between age and BMI, although this is unclear graphically.

Now we view the number of women to men who are smokers to nonsmokers:

```
# Convert to dataframe
df = as.data.frame(insurance)

# Create ggplot object
p = ggplot(insurance, aes(x=sex, fill=smoker))
p + geom_bar(position="stack")
```




We can see here that, as expected, there are slightly more smokers amongst men than there are smokers amongst women. Another important thing to note is that there are almost equally as many men represented in the dataset as there are women.

Let's see the actual ages represented in this dataset, as well as sex per age, and try to see if it's balanced:

```
p = ggplot(df, aes(x=age, fill=sex))
```

```
p + geom_bar(position="stack")
```



We can see here that there is such a high representation of ages under 20. To eliminate bias towards such younger age groups, it may help to remove all input from ages under 20.


Let's quickly view the number of smokers here:

```
insurance %>% group_by(smoker) %>% summarise("count"=n())
```

```
> insurance %>% group_by(smoker) %>% summarise("count"=n())
# A tibble: 2 x 2
  smoker count
  <chr>   <int>
1 no      1064
2 yes      274
> |
```

We've got 274 smokers in this dataset. Now let's look at this graphically:

```
ggplot(insurance, aes(x=age, y=charges, color=smoker)) + geom_point()
```




We can see that many of those smokers are the outliers here, in which they are charged much more for insurance than non-smokers.

### Data Cleaning:

As we saw from the data exploration part above, this dataset seems to be mostly straightforward with no NA's or missing data. Although, we did notice that there may be a bias towards ages younger than 20, so we will not remove all data points of those younger than 20.

Now we exclude rows using subset function with condition of including ages greater than or equal to 20, and see how this looks graphically:



We can see that it now looks much more balanced.



## Data Preprocessing:

Now, we will bin both age groups and charges which may be helpful when processing the data further. We will bin age groups into 5 different bins: "twenties", "thirties", "fourties", "fifties" and "sixtiesplus".

```
myinsurance <- myinsurance %>%
```

```
mutate(agegroup = cut(age, breaks=c(-Inf, 29, 39, 49, 59, Inf),labels=c("twenties",  
"thirties", "fourties", "fifties", "sixtiesplus")))
```

```
head(myinsurance)
```

```
> head(myinsurance)
  age  sex  bmi children smoker  region  charges agegroup
3  28 male 33.000         3    no southeast  4449.462 twenties
4  33 male 22.705         0    no northwest 21984.471 thirties
5  32 male 28.880         0    no northwest  3866.855 thirties
6  31 female 25.740        0    no southeast  3756.622 thirties
7  46 female 33.440        1    no southeast  8240.590 fourties
8  37 female 27.740        3    no northwest  7281.506 thirties
```

## Clustering:

Firstly we remove class labels and create dummy variables:

```
df = myinsurance
```

```
predictors <- df %>% select(-c(agegroup, region))
```

```
head(predictors)
```

```
> head(predictors)
  age  sex  bmi children smoker  charges
3  28 male 33.000         3    no  4449.462
4  33 male 22.705         0    no 21984.471
5  32 male 28.880         0    no  3866.855
6  31 female 25.740        0    no  3756.622
7  46 female 33.440        1    no  8240.590
8  37 female 27.740        3    no  7281.506
```

```
##create dummies
```

```
dummy <- dummyVars(charges ~ ., data = predictors)
```

```
dummies <- as.data.frame(predict(dummy, newdata = predictors))
```

```
##include charges
```

```
dummies$charges = myinsurance$charges
```

```
##rename predictors
```

```
predictors <- dummies
```

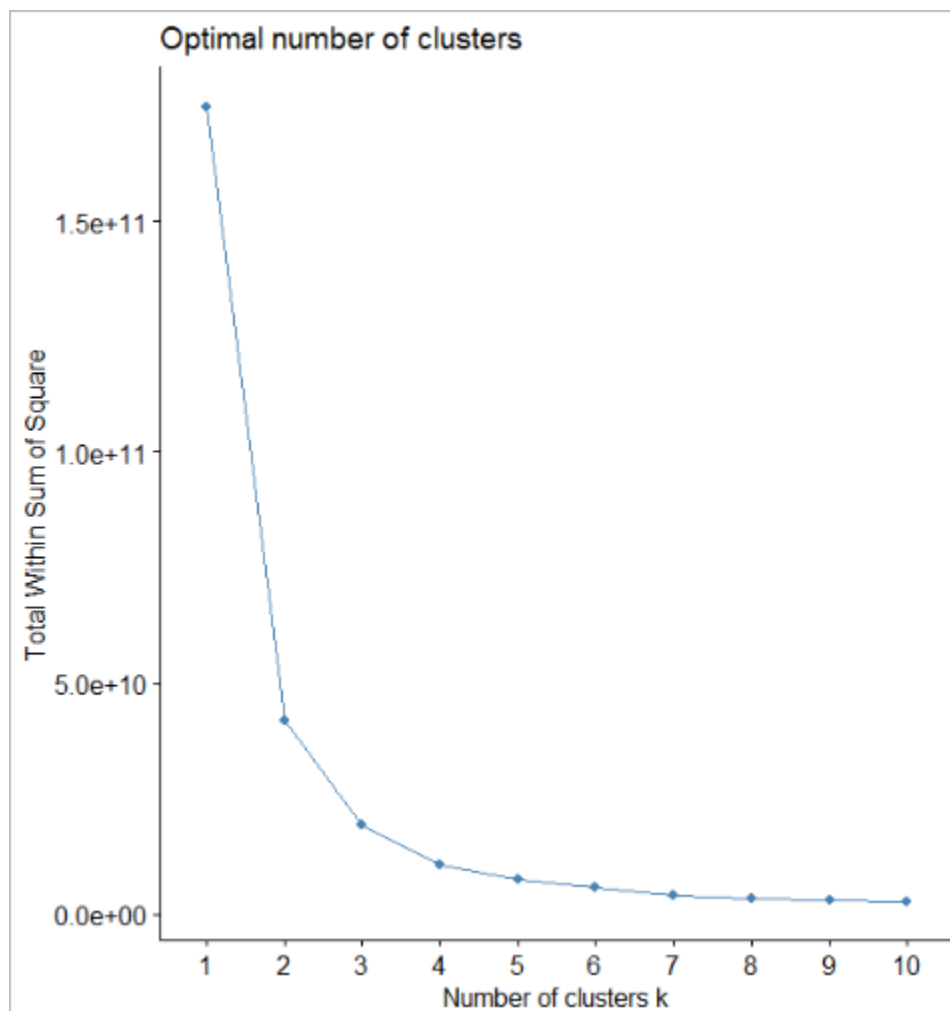
Now I will try K-means:

```
##set seed
```

```
set.seed(123)
```

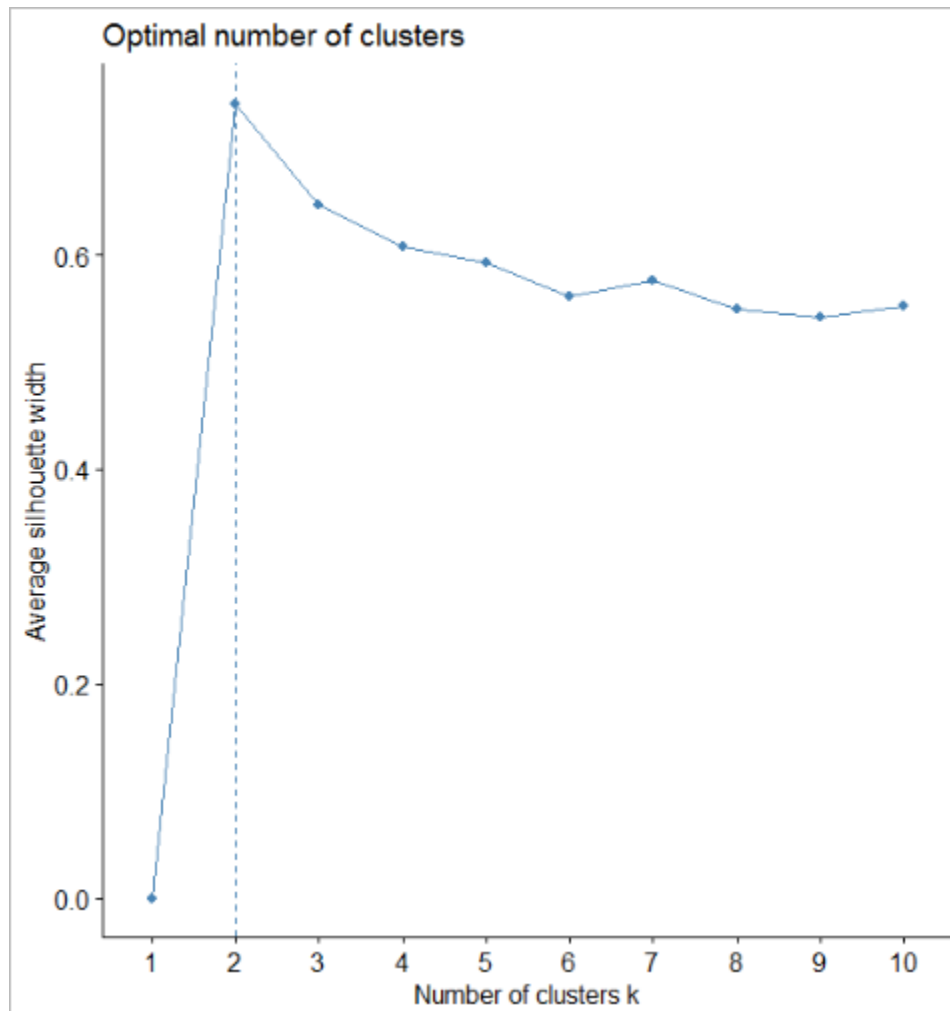
```
##find K; find the knee
```

```
fviz_nbclust(predictors, kmeans, method = "wss")
```



```
##use silhouette to find K
```

```
fviz_nbclust(predictors, kmeans, method = "silhouette")
```



The knee suggests a K of 4 but the silhouette score suggests K = 2. Using K = 4 as suggestion by the plots we can fit our model using the k-means function.

```
# Fit the data
```

```
fit <- kmeans(predictors, centers = 4, nstart = 25)
```

```
# Display the kmeans object information
```

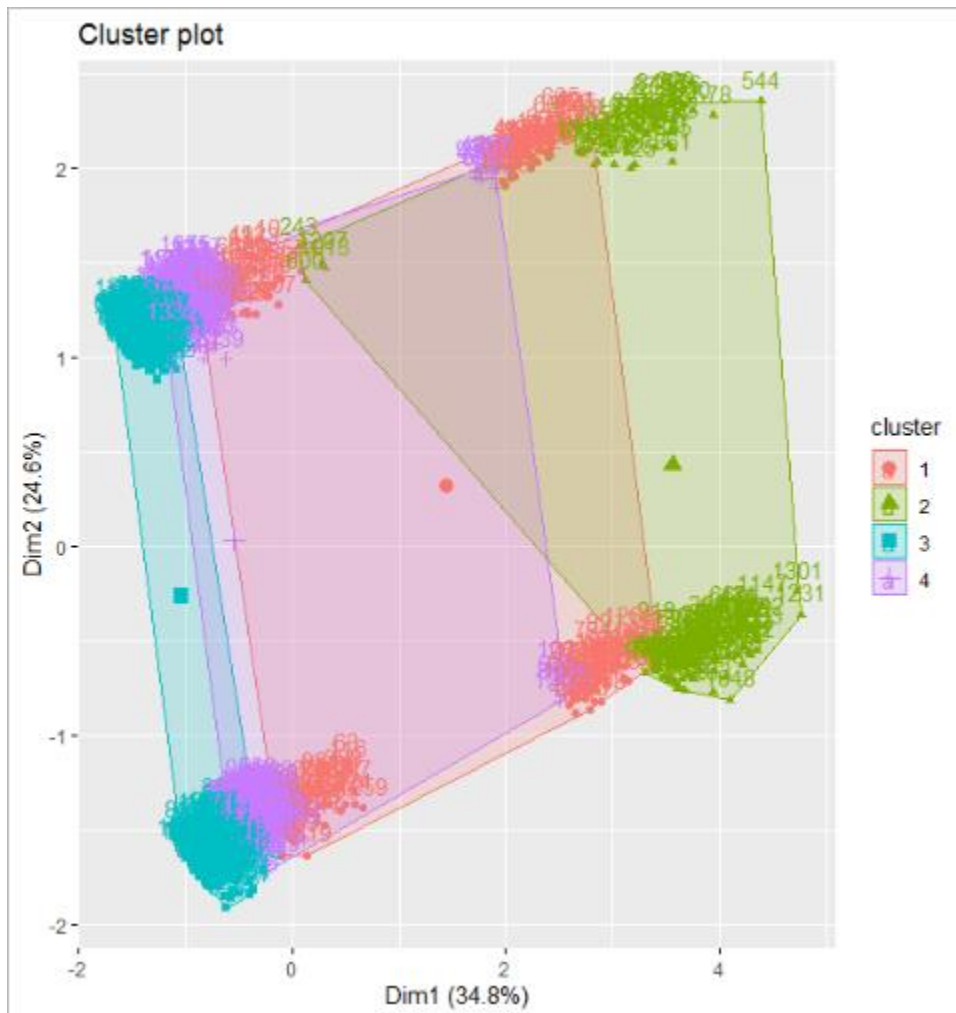
```
fit
```

```
K-means clustering with 4 clusters of sizes 164, 138, 481, 418
Cluster means:
  age sexfemale  sexmale    bmi children  smokerno  smokeryes  charges
1 42.78659 0.4756098 0.5243902 28.13838 1.250000 0.42682927 0.57317073 22850.673
2 41.94203 0.3695652 0.6304348 35.10152 1.231884 0.02898551 0.97101449 41963.079
3 31.94179 0.4989605 0.5010395 30.24236 1.143451 1.00000000 0.00000000 4803.633
4 52.04785 0.5430622 0.4569378 30.90012 1.148325 0.96172249 0.03827751 11375.730
```

Using the `fviz_cluster` function, we can visualize how the clusters are formed.

```
# Display the cluster plot
```

```
fviz_cluster(fit, data = predictors)
```



For comparison we can generate our own PCA plot and color the points based on their charges.

```
## Calculate PCA
```

```
pca = prcomp(predictors)
```

```
## Save as dataframe
```


```
rotated_data = as.data.frame(pca$x)
```

```
## Add original labels as a reference
```

```
rotated_data$Color <- df$charges
```

```
## Plot and color by labels
```

```
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Color)) +  
geom_point(alpha= 0.3)
```




The cluster plot can also be done on ggplot based on the cluster result from the K-means algorithm:

```
## Assign clusters as a new column
```

```
rotated_data$Clusters = as.factor(fit$cluster)
```

```
## Plot and color by labels
```

```
ggplot(data = rotated_data, aes(x = PC1, y = PC2, col = Clusters)) + geom_point()
```



## Classification

Now we use KNN to make predictions based on smokers:

```
set.seed(123)
```

```
ctrl = trainControl(method="cv", number = 10)
```

```
knnFit <- train(smoker ~ ., data = myinsurance, method = "knn", trControl = ctrl,  
preProcess = c("center","scale"))
```

*knnFit*

```
k-Nearest Neighbors

1201 samples
  7 predictor
  2 classes: 'no', 'yes'

Pre-processing: centered (12), scaled (12)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 1080, 1081, 1080, 1081, 1081, 1082, ...
Resampling results across tuning parameters:

   k  Accuracy   Kappa
   5  0.9075869  0.6734083
   7  0.9026076  0.6473205
   9  0.9018090  0.6354553

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 5.
```

From the output above, we can see that Accuracy and Kappa are reportedly high. Now, we'll attempt to control and find the best k value:

*set.seed(123)*

*ctrl = trainControl(method="cv", number = 10)*

*knnFit <- train(smoker ~ ., data = myinsurance, method = "knn", trControl = ctrl, preProcess = c("center","scale"), tuneLength = 15)*

*knnFit*

```
k-Nearest Neighbors


1201 samples
  7 predictor
  2 classes: 'no', 'yes'

Pre-processing: centered (12), scaled (12)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 1080, 1081, 1080, 1081, 1081, 1082, ...
Resampling results across tuning parameters:

   k  Accuracy   Kappa
   5  0.9075869  0.6734083
   7  0.9026076  0.6473205
   9  0.9018090  0.6354553
  11  0.8976699  0.6090430
  13  0.8960102  0.5968295
  15  0.8934962  0.5852356
  17  0.8860239  0.5465009
  19  0.8835237  0.5349892
  21  0.8835028  0.5346791
  23  0.8793431  0.5148305
  25  0.8768289  0.5018231
  27  0.8743288  0.4870325
  29  0.8726620  0.4771848
  31  0.8676758  0.4520441
  33  0.8676687  0.4511347

Accuracy was used to select the optimal model using the largest value.
The final value used for the model was k = 5.
```

*plot(knnFit)*



Accuracy varies between 0.97 and 0.98, which is quite good. However, to avoid overfitting, it's best to choose a higher K number as chances of overfitting decrease. Therefore, a K value of 18 may be best.

Now, we'll use grid search to find the best value of K:

*##setup a tuneGrid with the tuning parameters*

```
tuneGrid <- expand.grid(kmax = 3:7, kernel = c("rectangular", "cos"), distance = 1:3)
```

*## tune and fit the model with 10-fold cross validation, standardization, and our specialized tune grid*

```
kknn_fit <- train(smoker ~ ., data = myinsurance, method = 'knn', trControl = ctrl,  
preProcess = c('center', 'scale'), tuneGrid = tuneGrid)
```



*##Printing trained model provides report*

*kkn\_fit*

```
k-Nearest Neighbors
1201 samples
  7 predictor
  2 classes: 'no', 'yes'

Pre-processing: centered (12), scaled (12)
Resampling: Cross-validated (10 fold)
Summary of sample sizes: 1082, 1081, 1081, 1082, 1080, 1082, ...
Resampling results across tuning parameters:

  kmax  kernel  distance  Accuracy  Kappa
3      rectangular  1      0.9158848  0.7269340
3      rectangular  2      0.9142249  0.7232849
3      rectangular  3      0.9133778  0.7227479
3      cos         1      0.9158848  0.7269340
3      cos         2      0.9158501  0.7325850
3      cos         3      0.9116972  0.7176086
4      rectangular  1      0.9158848  0.7269340
4      rectangular  2      0.9142249  0.7232849
4      rectangular  3      0.9133778  0.7227479
4      cos         1      0.9150445  0.7238570
4      cos         2      0.9158501  0.7325850
4      cos         3      0.9133778  0.7209788
5      rectangular  1      0.9158848  0.7269340
5      rectangular  2      0.9142249  0.7232849
5      rectangular  3      0.9133778  0.7227479
5      cos         1      0.9150445  0.7238570
5      cos         2      0.9158501  0.7325850
5      cos         3      0.9142112  0.7225564
6      rectangular  1      0.9158848  0.7269340
6      rectangular  2      0.9142249  0.7232849
6      rectangular  3      0.9133778  0.7227479
6      cos         1      0.9150445  0.7238570
6      cos         2      0.9158501  0.7325850
6      cos         3      0.9142112  0.7225564
7      rectangular  1      0.9158848  0.7269340
7      rectangular  2      0.9142249  0.7232849
7      rectangular  3      0.9133778  0.7227479
7      cos         1      0.9150445  0.7238570
7      cos         2      0.9125167  0.7181108
7      cos         3      0.9133778  0.7188768

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were kmax = 7, distance = 1 and kernel = rectangular.
```

Now we can apply the model based on the data above:

*## Predict*

```
pred_knn <- predict(kkn_fit, myinsurance)
```

*## Generate confusion matrix*

```
myinsurance$smoker = as.factor(myinsurance$smoker)
```

```
confusionMatrix(myinsurance$smoker, pred_knn)
```

## Confusion Matrix and Statistics

	Reference	
Prediction	no	yes
no	957	0
yes	0	244

Accuracy : 1  
95% CI : (0.9969, 1)  
No Information Rate : 0.7968  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 1

McNemar's Test P-Value : NA

Sensitivity : 1.0000  
Specificity : 1.0000  
Pos Pred Value : 1.0000  
Neg Pred Value : 1.0000  
Prevalence : 0.7968  
Detection Rate : 0.7968  
Detection Prevalence : 0.7968  
Balanced Accuracy : 1.0000

'Positive' Class : no

## ## Result

```
knn_results = kknn_fit$results
```


```
knn_results <- knn_results %>%
```

```
  group_by(kmax, kernel) %>%
```

```
  mutate(avgacc = mean(Accuracy))
```

```
ggplot(knn_results, aes(x=kmax, y=avgacc, color=kernel)) +
```

```
  geom_point(size=3) + geom_line()
```



### Evaluation:

Using the better classifier kkn from previous part, we produce confusion matrix:

```
## Generate confusion matrix
```

```
myinsurance$smoker = as.factor(myinsurance$smoker)
```

```
cm = confusionMatrix(myinsurance$smoker, pred_knn)
```

## Confusion Matrix and Statistics

```

      Reference
Prediction no yes
no      957   0
yes      0 244

      Accuracy : 1
      95% CI : (0.9969, 1)
No Information Rate : 0.7968
P-Value [Acc > NIR] : < 2.2e-16

      Kappa : 1

McNemar's Test P-Value : NA

      Sensitivity : 1.0000
      Specificity : 1.0000
Pos Pred Value : 1.0000
Neg Pred Value : 1.0000
Prevalence : 0.7968
Detection Rate : 0.7968
Detection Prevalence : 0.7968
Balanced Accuracy : 1.0000

      'Positive' Class : no
```

Predict Class	Actual Class		
		Smoke = No	Smoke = Yes
	Smoke = No	957	0
	Smoke = Yes	0	244

*## Store the byClass object of confusion matrix as a dataframe*

*metrics <- as.data.frame(cm\$byClass)*

*## View the object*

*metrics*

```

cm$byClass
Sensitivity      1.000000
Specificity      1.000000
Pos Pred Value   1.000000
Neg Pred Value   1.000000
Precision         1.000000
Recall           1.000000
F1               1.000000
Prevalence       0.796836
Detection Rate   0.796836
Detection Prevalence 0.796836
Balanced Accuracy 1.000000

```

Now we calculate the precision and recall manually:

Predicted Class	Actual Class				
		Smoke = No	Smoke = Yes	Total	Recognition(%)
	Smoke = No	957	0	957	100%
	Smoke = Yes	0	244	244	100%
	Total	957	244	1201	
	Recognition(%)	100%	100%		


Finally we produce ROC plot:

```

## Get class probabilities for KNN
pred_prob <- predict(kknn_fit, myinsurance, type = "prob")
head(pred_prob)

## And now we can create an ROC curve for our model.
roc_obj <- roc((myinsurance$smoker), pred_prob[,1])
plot(roc_obj, print.auc=TRUE)

```



We can see that these performance measures makes our classifier look the same as accuracy.

### **Report:**

The insurance.csv dataset contains 1338 observations (rows) and 7 features (columns). The dataset contains 4 numerical features (age, bmi, children and expenses) and 3 nominal features (sex, smoker and region) that were converted into factors with numerical value designated for each level.

The purposes of this dataset are to look into different features to observe their relationship, from that plot a multiple linear regression based on several features of individual such as age, physical/family condition and location against their existing medical expense to be used for predicting future medical expenses of individuals that help medical insurance to make decision on charging the premium.

After exploring and cleaning data, using clustering and classification help us to predict and evaluate data. The model has shown good performance with 100% sensitivity. The

hyperparameter tuning was done at the end to check the difference and have seen minor difference in the model performance. The 'smoker' variable is one of the variables which play the most important in the model that did not include the age variable and the model that included the age variable has shown that the 'age' variable was most important.

The confusion matrix has shown the various performance metrics for both knn and kkn classifier. The comparison reveals that Highest Sensitivity and Highest Accuracy both were seen in the kkn classifier.

The kkn classifier has great accuracy and sensitivity, and decreased possible confusion, it can be concluded that the exercise has ended with a better performing model in comparison to the models created. The model was successful in fulfilling the goals of the project. It will make a decision safe enough for the organizations to predict and organize their appointments as necessary. It would be interesting to see the results with a decision tree model or even clustering techniques in case a re-visit is planned with this dataset.